

GRID CONTROL — Build Plan + As-Built Status

Safe word: **GRIDLOCK-RESUME-18JUNE**. Master sequence (Gaurav, unchanged): **Backend → Frontend → Stitch → Test brand → Launch**. Plan written Jun 18 2026 · As-built status updated Jun 20 2026.

STATUS: Backend (Master-Sequence Step 1) = ✔ COMPLETE. Phases 1–5 shipped + validated. Total Phase-5 validation spend ≈ **\$1.83** (budget was \$15). Now moving to **Step 2 — Front-end**. Legend: ✔ shipped · ⚠ shipped with deviation · 🛑 not done (reason) · ⌚ pending.

0 · Why this exists

After the Jun 18 deep-fetch of ~71 repos (`REPO_RESEARCH_INDEX.md` §C), two system-level pieces — **semantic memory** and **observability/cost** — touch every agent, so they went FIRST, before per-agent expertise packs. **caveman** is separate: a Claude-Code session-ergonomics skill (cuts build-chat tokens), NOT a production-cost lever. Picsart Agents (reviewed Jun 18) added one backlog item only: WhatsApp/Telegram as future Chief-of-Staff channels (post-launch).

1 · Phase 1 — Foundation upgrades ✔ COMPLETE

#	Item	As-built status
1a	Semantic memory backbone	⚠ Shipped as Voyage <code>voyage-3-lite</code> (512-dim) → Supabase pgvector, NOT the <code>mem0ai</code> lib (chose direct embed→pgvector for control + simplicity). Migrations 007/008. <code>agents/_lib/_mem0_client.py</code> . 2 scopes (<code>grid_control</code> + <code>brand:<slug></code>), not 3 — Gaurav's call, keeps agent context clean. strategy-agent wired as first consumer (recall before, write winner after).
1b	Langfuse observability	⚠ Langfuse Cloud (free tier), not self-hosted. <code>cost_reporter</code> KEPT (not retired) — <code>utils/pricing.py</code> is the cost source of truth; Langfuse is additive. Per-generation cost attribution fixed (<code>start_generation</code> context mgr + <code>base_agent.lf_generation</code>); strategy-agent wired + verified on cloud. Other agents adopt the one-line wrap when next touched.
1c	rate_limit decorator-order bug	✔ Already fixed earlier (commit 32aa062, Jun 16) — <code>@bp.route</code> outer / <code>@rate_limit</code> inner on all 10 endpoints. Audit-confirmed.
1d	AgentShield baseline	✔ Baseline saved (<code>docs/AGENTSHIELD_BASELINE_18JUNE_FIXED.txt</code>). Grade C (62/100); 2 CRITICALs eliminated (sudo allow-rules removed + <code>Bash(sudo *)</code> denylisted). Re-run before first client.
1e	caveman skill	✔ Installed opt-in only (<code>/caveman ...</code>), not always-on (CLAUDE.md already caps length). Dev-ergonomics, doesn't touch the 18 agents.

2 · Phase 2 — Data & research hardening ✔ COMPLETE

#	Item	As-built status
2a	Scrapling	✔ <code>agents/_lib/_scrapling_client.py</code> (Fetcher + stealthy_headers). New <code>trend_researcher.scrape_competitor_websites()</code> reads <code>brand_profile.competitor_websites</code> . Additive — NOT an Apify replacement (Apify stays canonical for IG/YT, fixed-cost Starter plan).
2b	Supadata video→transcript	✔ <code>agents/_lib/_supadata_client.py</code> . <code>trend_researcher._extract_whisper_transcripts</code> now Supadata-first / Whisper-fallback. Real fix: yt_dlp was silently missing → transcripts broken for months.
2c	Google/Meta/GA4 metrics MCP	🛑 Not installed. Native Meta Ads/IG MCP already connected (~80 tools cover Meta). Ryze unified MCP redundant + blocked (signup needs a website that doesn't exist yet). GA4 + Google Ads metrics deferred until those upstream accounts actually exist — not blocking.

3 · Phase 3 — Expertise packs ✔ COMPLETE (lighter than original spec)

⚠ Shipped as prompt-only expertise-pack tables in each `managed_agents/prompts/*.md` (sub-task → skill mapping before the AutoResearch loop). **13 agents wired, 1 N/A** (carousel-designer = deterministic renderer). No per-agent `program.md` files (the original exit criterion) — the prompt-table approach was lighter and sufficient. Packs used: `arnabbagxd/Brand-building-skills` , `phuryrn/pm-skills` , `carlosarthurrl/business-playbook` , `AgriciDaniel/claude-ads` + existing higgsfield/geo/searchfit/design/data/marketing/sales packs.

Roster correction: the Jun-18 draft listed a "Lead Generator" (3o) and "AI-setter" (3p) as agents. Those are **NOT in the locked 18-agent roster**. No new agents were created — `dm-customer-hunter` absorbed local-biz prospecting instead (see 4b). The 18-agent roster is unchanged.

4 · Phase 4 — Routing policy + new capabilities COMPLETE

#	Item	As-built status
4a	Cheaper tier for high-volume grunt	 Haiku GRUNT tier, SPLIT not swap. <code>model_gateway.grunt_model()</code> = <code>claude-haiku-4-5</code> · <code>community-manager</code> + <code>dm-customer-hunter</code> split 2-stage: Haiku triages ALL items (categorize / ICP-score, incl. spam/off-ICP) → Sonnet writes brand-voice only for winners . Honors the Jun-9 Sonnet-floor lock (brand-facing text never on Haiku). Separate cost lines. (No "Lead Generator"/"AI-setter" — those agents don't exist.)
4b	Local-biz prospecting	 Apify Compass <code>crawler-google-places</code> , NOT <code>omkarcloud/google-maps-scraper</code> (botasaurus/Chromium = heavy on Railway). Reuses existing <code>_apify_run_sync</code> · <code>discover_via_google_maps()</code> gated on <code>brand_profile.gmaps_queries</code> (off by default, brand-agnostic). Flows through the same Haiku triage. Zero new deps.
4c	Google Ads API client	  <code>agents/_lib/_google_ads_client.py</code> wired. OAuth chain verified LIVE (returns MCC <code>customers/7023267019</code>). Dev token in TEST access; Basic Access application submitted Jun 19 → Google review 5–14 business days . Live campaign data unlocks on approval.

5 · Phase 5 — Validation gate COMPLETE (≈ \$1.83 of \$15 budget)

#	Item	As-built status
5a	F4 magnet copy via funnel_specialist	 Ran for askgauravai → winner Variant C, in <code>pending_approval/funnel-specialist/</code> · \$0.115 . Fixed a real bug: <code>run_autoresearch_loop</code> <code>max_tokens</code> 6000→16000 (3-variant JSON was truncating).
5b	Paid-ops validation	 Tested DRY (\$0, NO real ads) . Gaurav vetoed ad spend — never discussed + not in the master sequence yet (paid = Step 5, post brand-onboard). Circuit-breaker verified on all 4 gate states (switch-off→block, on+under-cap→allow, on+est>cap→block, on+ledger-over-cap→block). Real-spend ledger confirmed live.
5c	End-to-end smoke test	 Lean keystone run (not all 13 agents — most unchanged except prompt packs). Ran strategy-agent (Opus, \$1.554 — 3 passes, Rule-10 forced 2 citation re-prompts → 13/13) + community-manager + dm-customer-hunter (Haiku split, ~\$0.04) on synthetic inputs (since deleted). Proven live: semantic recall+write · Rule-10 provenance loop · Haiku triage→Sonnet split · cost capture · Langfuse traces in cloud.

Findings (non-blocking): (1) Langfuse trace-level cost was \$0 — fixed via `start_generation` wrap (strategy done; others adopt when touched). (2) Notion push broken everywhere — `NOTION_PAGE_ID` is a URL not a UUID + Notion bill unpaid; outputs save locally to `pending_approval/`.

Exit criteria of Phase 5 = exit of Master-Sequence Step 1 (Backend).  **MET.**

6 · Cost-control actions (Jun 19–20)

Gaurav recharged Anthropic credits = **TEST USE ONLY; no automation runs without an explicit green signal**. -  Railway **scheduler-worker service REMOVED** (was firing askgauravai 06:30 + offgrid 07:00 IST daily — Apify+Claude spend). `scheduler/schedule_config.json` jobs also set `enabled=false`. -  Local crontab already disabled (Jun 16). Paid-ops **kill-switch OFF (safe)**; `GRID_PAID_OPS` never persisted. -  A separate `Trading-Bot` paper loop (`run_v2_paper.py` , until Jun 30) uses **no** Anthropic credits — left running.

7 · Housekeeping (Jun 20)

-  **Dead code removed:** `jarvis_config/` , `jarvis-skills/` , `scripts/install_jarvis.sh` (April voice experiment, never imported) + 4 stale AgentShield-flagged allow-rules. Live `/api/jarvis/query` route kept.
-  Orphan candidate flagged: `main.py` (stale launcher, superseded by `dashboard_api.py`) — awaiting decision.
-  **Backend committed + pushed** to `gridcontrol-rebuild` (7 commits): scheduler pause, 4a/4b+funnel-fix, Mem0+Langfuse+cost-fix, Scrapling+Supadata, Phase-3 packs, Google Ads client, docs/config, jarvis cleanup.

8 · NEXT — Master-Sequence Step 2: Front-end

Build the **GRID CONTROL** owner-facing command deck. Direction locked Jun 20 (`docs/FE_MASTER_PROMPT.md`):

- **Mental model:** Owner → **Chief of Staff** (the one persona they talk to) → **the Team on the factory floor** (avatar'd specialist *characters* who visibly work and bring finished pieces up for approval). Like an owner directing a manager and watching the team build it.
- **THE SECRET (hard rule):** the client UI **never** exposes backend machinery — no model names, costs, tokens, JSON, provenance, agent slugs, or infra. The team are *people*; the technology is invisible. Cost/token transparency is **internal ops only**, removed from the client UI.
- **Character/avatar system** is a design pillar (warm mascot energy, one per role).
- Design vibe to be chosen from Gaurav's **reference links** (slot in FE prompt \$7 + \$13).
- Build path: design against mock data → wire to the live API (`web-production-175d5...`) → translate every technical field into character/outcome language at the boundary.

9 · Deferred (not in this plan)

Higgsfield PRO + per-agent creative work (Step 5) · ManyChat F3 / live AI-setter (needs live offer + 2 links) · Open Design mining · WhatsApp/Telegram CoS channels (post-launch) · Postiz on Hetzner (post-launch) · GA4/Google-Ads live metrics (when those accounts exist) · roll Langfuse cost-wrap out to remaining agents.

Related: [FE MASTER PROMPT](#) · [REPO RESEARCH INDEX](#) · [ONBOARDING CLONE SPEC](#) · [CREATIVE ENGINE REBUILD](#) · [AGENT MODEL FIT RESEARCH](#).