

Grid Control — Working Process

Living document. Source of truth for how Grid Control (GC) works and what does what.

Last updated: 2026-06-25 · Status: WIRE phase (Stage 5 of 10).

1. What Grid Control Is

Grid Control is a client-ready, multi-tenant AI marketing operating system. A brand owner (the client) signs up, onboards their brand once, and an orchestrated team of AI agents researches, plans, creates, and (on approval) publishes their marketing — with a human approval gate on every output and a per-brand cost meter behind it.

- **Client app** = chat-console-primary. Atlas (the lead agent) guides the user; work appears inline in chat as approve-cards, with a Live Work Feed.
- **Real-data mandate / zero assumptions:** every data point traces to a real scrape, API call, or user input. Never fabricated. If a fact can't be verified, the system stops and asks.

2. Identities — who is who

Identity	Role	Notes
grid.admin1@gmail.com	GC platform owner / developer	Owns the Google Cloud project, the OAuth app, and all GC Google infra. The landlord. Never a tenant/brand.
Super-admin app user	GC operator	Currently offgridcreativesai@gmail.com (profiles.is_super_admin=true). Sees all brands; runs operator mode.
Client / tenant	A brand owner	Self-signs-up (email or Google). Sees only their own brand(s). Each brand fully isolated under brands/<slug>/.

Test brands (run through the real client flow): **Third Gen Tribe** (first), then AskGauravAI, OffgridCreatives.

3. The Client Lifecycle (the core flow)

Sign in → Onboarding form (Atlas) → Connect accounts → Pre-flight validation
→ Research report (paid) → Review / approve / 1 change
→ Foundation + Memory seed → First task handoff → Ongoing agent ops

3.1 Sign in

Email + password, magic link, or **Google** (enabled). New user with no brand → routed to onboarding wizard. Returning user with a brand → straight to Command Center. (Gated on live /api/brands, not localStorage.)

3.2 Onboarding form — Atlas-guided

7 screens: (0) new vs existing → (1) basics → (2) audience & goal → (3) voice & safety → (4) platforms, competitors, identity → (5) review → **(6) connect accounts** *(to build)*. Submits to POST /api/auth/create-brand → Supabase brands + brand_members(admin) + filesystem brand_profile.json.

3.3 Connect accounts (new 7th screen)

Per-platform connect buttons (Instagram/LinkedIn/YouTube/X), Noimos-style. Each connection reflects on the Connections page (same per-brand connection system). **Skippable**. Connected → report uses first-party insights (onboarding_connected); skipped → public research only (cold_sellable).

*Constraint: IG/LinkedIn insight + posting scopes need platform App Review (Phase B). Build the UI now, light up platforms as approved.

3.4 Pre-flight validation gate (cheap — no Opus)

Before any paid run, resolve every handle in real time (own IG + each competitor) via a light public-profile lookup. Any handle that doesn't resolve → Atlas stops in chat, asks the user to confirm/correct → re-check → only then start the report. **The system never assumes a handle is real.**

3.5 Research report (the onboarding deliverable — PAID)

brand-book/generate → agents/brand_book_v7 (Opus + real research/scraping + IG insights). Runs async. Atlas: "generating..." → "ready". This is the brand audit / Foundation, not raw agent dumps.

3.6 Review, approve, or change (ONE paid change cycle)

In-chat approve-card: short summary + **Open full PDF** · **Approve** · **Request changes**.

- **Approve** → Foundation locked.
- **Request changes** → captured as change_request → **scoped re-research only** (re-run just the changed section, reuse cached research), shown once. **Hard cap: 1 paid revision** (revision_count enforced in code).
- **Further changes** → done later in the Memory section (free, no model calls).
- **Brand colors**: if left empty / "you decide" → system proposes palette options → user approves one → locked.

3.7 Approve → Foundation + Memory seed

brand-book/approve writes Foundation → brand_profile.json + voice_profile.json + brand_narrative, and **seeds editable Memory entries** (positioning, audience, voice, competitors, do-not-say). This is the final source-of-truth agents read.

3.8 First-task handoff

Atlas proposes the first task from the brand's bottleneck (e.g. Third Gen Tribe inactive 15 months → "revival week" → Content Planner).

4. Cost Model (per brand — feeds GC pricing)

Pre-flight checks	→ cheap	(light lookups, no model)
First full report	→ PAID	(Opus + research) ← the main onboarding cost
≤ 1 scoped revision	→ cheap	(only changed section, cached research reused)
Memory edits	→ free	(no model calls)
Ongoing agent runs	→ metered	(every run logged per brand)

Tracking already built: agents/_lib/cost_reporter.py logs full run spend (Anthropic + Apify + FAL) → Supabase usage_logs + agent_runs. /api/billing/usage + /api/brands/<slug>/costs aggregate per-brand monthly cost + by-agent. To surface on the System page.

5. The Agent Roster — what does what

#	Agent	Role
0	CEO Brain	Orchestrator, session state, Notion push
1	Strategy Agent	90-day roadmap + competitor research (real scrapes)
2	Content Planner	30-day calendar
3	Script Writer	Scripts/hooks/captions + voice check
4	Creative Director	AI video/image (runs after founder records)
5	Ad Strategist	Paid funnels (only when budget confirmed)
6	Data Analyst	Real-account metrics, weekly score
7	Funnel Specialist	Conversion journey
8	Trend Researcher	Apify + Whisper + clustering
9	Website Agent	Site + GA4 + Search Console
10	Brand Guardian	SOUL/brand-safety check on all content
11	SEO/AEO Agent	Search + answer-engine optimisation
12	Email Marketing	Email sequences
13	Community Manager	Replies, mentions
14	DM Customer Hunter	Inbound + warm DMs
15	Carousel Designer	7-slide carousel → editorial HTML/CSS
16	Trend Sentinel	Daily PIVOT/TRACK/STAY (pure math)
17	Performance Tracker	Winning/dead patterns (pure math)

Class-1 (decision) agents = pure math, no model. Class-2 (generation) = Claude, every output cites data_provenance (source file + path + value).

6. Approval Gate & Publishing

Every agent output → brands/<slug>/outputs/pending_approval/. Human approves → outputs/approved/. Nothing auto-publishes. Publishing per platform (IG/LinkedIn/YouTube Shorts → automation; X → manual upload) only after explicit "approved".

7. Memory System

Per-brand memory (Supabase brand_memory + the Memory page). Seeded from the approved report; user/agents can add/edit entries over time. Becomes the evolving source of truth alongside brand_profile.json + voice_profile.json.

8. Tech & Infra

- **Frontend:** React 19 + Vite + Tailwind v4 + shadcn + TanStack Query v5 + Zustand + react-router v7. Dev on fixed port **localhost:5280** (strictPort; not 5173).
 - **Backend:** Flask `/api/*` on 5001, `/api` proxy via Vite. Blueprints in `routes/*.py` (brands, agents, content, brain, billing, connections, system, leads).
 - **Auth:** Supabase (project `mnivbrehrwgndxcgega`). JWT (`Authorization: Bearer`) or legacy X-Dashboard-Secret operator bypass. Google OAuth enabled (basic scopes).
 - **Data:** Supabase (brands, brand_members, usage_logs, agent_runs, brand_memory, ...). Per-brand secrets in `brands/<slug>/ .env` (gitignored).
 - **Deploy:** Vercel (dashboard) + Railway (Flask API + scheduler worker). Custom domain pending.
-

9. Current Build Status (WIRE — Stage 5 of 10)

Done: portal reset to empty · Google sign-in enabled · unique port 5280 · onboarding form → brand creation verified (Third Gen Tribe).

Next (agreed order): (1) fix data + form NA/comma bug → (2) live handle-validation gate → (3) report generate + Atlas status loop → (4) review/approve/one-change card → (5) approve → Foundation + Memory seed → (6) first-task handoff. Connect screen = 7th step in parallel.

10. Open Items / Risk Flags

- **Platform App Review** (Meta/LinkedIn insight + posting scopes) before reliable connect/publish.
- **Legal before paying clients:** ToS + Privacy Policy + DPA; encrypt stored tokens at rest; restrict CORS; confirm operator mode never client-reachable.
- **Cost control:** in-flight lock, one scheduler, budget caps (post the Jun credit-drain incident).
- Surface per-brand cost on the System page.